

MongoDB 社区版 on GKE部署说明

1. GKE集群安装

利用gcloud命令行工具一键创建GKE集群，在创建集群时，可通过“--cluster-dns”参数启用cloud dns，并通过--cluster-dns-scope参数将范围设置为VPC。启用该参数，可通过cloud dns维护kubernetes内pod与service的域名记录，而非Kubernetes默认的core-dns/kube-dns。启用该参数的好处时，可以让位于同一VPC内的GCE实例或者不同GKE集群之间的pod互相解析域名，实现服务发现机制的统一管理。

同时启用workload identity: 做velero 备份

[Using Cloud DNS for GKE | Kubernetes Engine Documentation](#)

```
gcloud container clusters create <cluster-name> \
--cluster-dns clouddns --cluster-dns-scope vpc \
--cluster-dns-domain <example.com> \
--zone=us-central1-a \
--machine-type=n2-standard-8 \
--num-nodes=3 \
--workload-pool=<PROJECT_ID>.svc.id.goog
```

The screenshot shows the Google Cloud Platform Cloud DNS interface. On the left, there's a sidebar with navigation links: Load balancing, Cloud DNS (which is selected and highlighted in blue), Cloud CDN, Cloud NAT, Traffic Director, Service Directory, Cloud Domains, and Private Service Connect. Below the sidebar, there are sections for Marketplace and Release Notes.

The main area is titled "RECORD SETS" and "IN USE BY". It includes buttons for "ADD RECORD SET", "DELETE RECORD SETS", and "REFRESH". A "Filter" input field is present. The table lists various record sets with columns for "DNS name", "Type", "TTL (seconds)", and "Routing policy". Most entries have a TTL of 30 seconds and a Default routing policy. Some entries like "external.mongo." and "external.mongo." have higher TTL values (21600). The "Type" column shows mostly SRV and A records, with one SOA record.

DNS name	Type	TTL (seconds)	Routing policy
_dns-tcp._tcp.kube-dns.kube-system.svc.external.mongo.	SRV	30	Default
_dns._udp.kube-dns.kube-system.svc.external.mongo.	SRV	30	Default
_http._tcp.default-http-backend.kube-system.svc.external.mongo.	SRV	30	Default
_https._tcp.kubernetes.default.svc.external.mongo.	SRV	30	Default
_mongodb._tcp.mongodb-small-cluster.svc.mongodb.svc.external.mongo.	SRV	30	Default
default-http-backend.kube-system.svc.external.mongo.	A	30	Default
dns-version.external.mongo.	TXT	30	Default
external.mongo.	NS	21600	Default
external.mongo.	SOA	21600	Default
kube-dns.kube-system.svc.external.mongo.	A	30	Default
kubernetes.default.svc.external.mongo.	A	30	Default
metrics-server.kube-system.svc.external.mongo.	A	30	Default
mongodb-small-cluster-0.mongodb-small-cluster.svc.mongodb.svc.external.mongo.	A	30	Default
mongodb-small-cluster-1.mongodb-small-cluster.svc.mongodb.svc.external.mongo.	A	30	Default
mongodb-small-cluster-2.mongodb-small-cluster.svc.mongodb.svc.external.mongo.	A	30	Default
mongodb-small-cluster-svc.mongodb.svc.external.mongo.	A	30	Default

Cloud DNS统一管理GKE集群内服务域名记录

使用PD-SSD作为默认的storage class, 并支持storage expansion:

```
kubectl patch sc standard -p
'{"metadata": {"annotations": {"storageclass.kubernetes.io/is-default-class": "false"} }}'

kubectl patch sc premium-rwo -p
'{"metadata": {"annotations": {"storageclass.kubernetes.io/is-default-class": "true"} }}'
```

2. MongoDB Operator部署

2.1 克隆MongoDB Operator代码至本地，并切换至MongoDB Operator目录

```
git clone https://github.com/mongodb/mongodb-kubernetes-operator.git
```

2.2 创建CRD资源，并查看crd资源的创建

```
# 创建CRD资源
kubectl apply -f
config/crd/bases/mongodbcommunity.mongodb.com_mongodbcommunity.yaml
查看CRD资源的创建
kubectl get crd/mongodbcommunity.mongodbcommunity.mongodb.com
```

2.3 安装role与role-bindings，并查看资源的创建

```
# 创建role与role-bindings
kubectl apply -k config/rbac/ --namespace <my-namespace>
# 查看资源的创建
kubectl get role mongodb-kubernetes-operator --namespace <my-namespace>
kubectl get rolebinding mongodb-kubernetes-operator --namespace <my-namespace>
kubectl get serviceaccount mongodb-kubernetes-operator --namespace <my-namespace>
```

编辑“`config/manager/manager.yaml`”配置文件，添加`CLUSTER_DOMAIN`环境变量，将环境变量的值设置为创建GKE集群时“`cluster-dns-domain`”指定的域名。

```
env:
- name: WATCH_NAMESPACE
  valueFrom:
    fieldRef:
      fieldPath: metadata.namespace
- name: POD_NAME
  valueFrom:
    fieldRef:
      fieldPath: metadata.name
- name: OPERATOR_NAME
  value: mongodb-kubernetes-operator
- name: AGENT_IMAGE
  value: quay.io/mongodb/mongodb-agent:11.12.0.7388-1
- name: VERSION_UPGRADE_HOOK_IMAGE
  value: quay.io/mongodb/mongodb-kubernetes-operator-version-upgrade-post-start-hook:1.0.4
- name: READINESS_PROBE_IMAGE
  value: quay.io/mongodb/mongodb-kubernetes-readinessprobe:1.0.8
- name: MONGODB_IMAGE
  value: mongo
- name: MONGODB_REPO_URL
  value: docker.io
- name: CLUSTER_DOMAIN
  value: <example.com>
```

编辑operator配置文件

2.4 安装Operator

```
#安装Operator
kubectl create -f config/manager/manager.yaml --namespace <my-namespace>
#查看Operator创建
kubectl get pods --namespace <my-namespace>
```

3. MongoDB副本集部署

3.1 MongoDB副本集部署建议

2.1.1 配置建议

	小规模	中等规模	较大规模
CPU/Memory/Disk	2C/8GB/512GB SSD	4C/16GB/1TB SSD	8C/32G/2TB SSD
开源版本	4.4	4.4	4.4
副本数量	3	3	3

2.1.2 部署多套副本集建议

	部署
开发/测试	与应用开发测试相同的Cluster不同Namespace
生产	单独的Cluster

3.2 MongoDB配置

将“config/samples/mongodb.com_v1_mongodbcommunity_cr.yaml”文件中的“<your-password-here>”替换成自己要设置的密码。

3.3 安装MongoDB副本集

```
kubectl apply -f config/samples/mongodb.com_v1_mongodbcommunity_cr.yaml --namespace <my-namespace>
```

3.4 查看MongoDB副本集部署状态

```
kubectl get mongodbcommunity --namespace <my-namespace>
```

3.5 获取MongoDB副本集连接url

```
kubectl get secret <metadata.name>-<auth-db>-<username> -n <my-namespace> \
-o json | jq -r '.data | with_entries(.value |= @base64d)'
```

<metadata.name>, <auth-db>以及<username>的信息参照下表替换成对应的内容。

Variable	Description	Value in Sample
<metadata.name>	Name of the MongoDB database resource.	example-mongodb
<auth-db>	Authentication database where you defined the database user.	admin
<username>	Username of the database user.	my-user

输出内容类似如下，利用输出的字符串连接集群

```
{
  "connectionString.standard": "mongodb://<username>:<password>@example-mongodb-0.example-mongodb-svc.mongodb.svc.cluster.local:27017,example-mongodb-1.example-mongodb-svc.mongodb.svc.cluster.local:27017,example-mongodb-2.example-mongodb-svc.mongodb.svc.cluster.local:27017/admin?ssl=true",
  "connectionString.standardSrv": "mongodb+srv://<username>:<password>@example-mongodb-svc.mongodb.svc.cluster.local/admin?ssl=true",
  "password": "<password>",
  "username": "<username>"
}
```

3.6 可通过调整MongoDBCommunity内members的数量实现集群的扩缩容，通过调整version来实现对MongoDB集群版本的管理。

```
apiVersion: mongodbcommunity.mongodb.com/v1
kind: MongoDBCommunity
metadata:
  name: example-mongodb
spec:
  members: 3
  type: ReplicaSet
  version: "4.2.7"
  security:
    tls:
```

```
enabled: true
certificateKeySecretRef:
    name: <tls-secret-name>
caConfigMapRef:
    name: <tls-ca-configmap-name>
```

4. MongoDB pod容量调整

通过查阅MongoDB Operator配置文档，可知MongoDB通过statefulset机制来对Pod的生命周期进行管理。因此Pod容量调整可通过设置MongoDB CRD描述中Statefulset的部分来实现。

```
statefulSet:
  spec:
    template:
      spec:
        # resources can be specified by applying an override
        # per container name.
        containers:
          - name: mongod
            resources:
              limits:
                cpu: "0.2"
                memory: 250M
              requests:
                cpu: "0.2"
                memory: 200M
          - name: mongodb-agent
            resources:
              limits:
                cpu: "0.2"
                memory: 250M
              requests:
                cpu: "0.2"
                memory: 200M
        ...
        ...
        ...
      
```

MongoDB CRD statefulset 部分描述

3.1、调整MongoDB 节点CPU与内存限制

在默认配置文件下安装MongoDB集群，Operator会为Master以及每个Replicaset节点设置CPU/Mem 的limit为0.2/250M，设置每个Replicaset节点的requests 0.2/200M。通过查寻其中一个Pod的描述，也可以验证此限制。

```
kubectl describe pod example-mongodb-0 -n mongodb
```

```
Limits:
  cpu:     1
  memory: 500M
Requests:
  cpu:    500m
  memory: 400M
```

接下来，将MongoDB的CPU/Mem limit上调至2/2048M，将CPU/Mem Request上调至1/1048M，应用配置文件，查看资源调整情况。

```
statefulSet:
  spec:
    template:
      spec:
        # resources can be specified by applying an override
        # per container name.
        containers:
          - name: mongod
            resources:
              limits:
                cpu: "2"
                memory: 2048M
              requests:
                cpu: "1"
                memory: 1024M
```

MongoDB 配置文件

```
kubectl apply -f <mongodb-conf>
```

可参考这个[配置文件](#)

重新应用配置文件后，MongoDB Operator会触发statefuset重新启动pod。

```
^Cadmin_zzcao_altostrat_com@kubectl:~/mongodb-kubernetes-operator/config/samples$ kubectl get pod -n mongodb -w
NAME                           READY   STATUS    RESTARTS   AGE
mongodb-kubernetes-operator-774b59c94f-nn5x1   1/1     Running   0          4h4m
mongodb-specify-pod-resources-0                 2/2     Running   0          160m
mongodb-specify-pod-resources-1                 2/2     Running   0          159m
mongodb-specify-pod-resources-2                 0/2     Init:0/2  0          8s
```

自动重启MongoDB Pod

```
Limits:
  cpu:     2
  memory: 2048M
Requests:
  cpu:    1
  memory: 1024M
```

Pod配置调整生效

3.2、调整MongoDB 存储卷

MongoDB Operator通过创建PVC实现数据持久化存储，Operator会为每一个MongoDB节点分配两个PVC存储卷，一个存储卷用于存储数据(默认10G)，一个存储卷用于存储日志(默认2G)。

调整MongoDB存储卷的过程分为三步：

- 更改MongoDB Operator CRD定义文件，修改存储容量，并重新应用该配置

```
kubectl apply -f config/samples/mongodb.com_v1_mongodbcommunity_cr.yaml
```

- 删除statefulset，MongoDB Operator会自动创建新的statefulset。

```
kubectl delete statefulset example-mongodb
```

- 更改pvc存储卷容量

```
kubectl patch pvc data-volume-example-mongodb-2 -p '{"spec": {"resources":{"requests":{"storage": "18Gi"}}}}'
```

```
kubectl get pvc
```

```
I have no name!@mongodb-specify-pod-resources-0:/ $ df -h
Filesystem      Size  Used Avail Use% Mounted on
overlay         95G   6.2G  89G  7% /
tmpfs          64M     0   64M  0% /dev
tmpfs          15G     0   15G  0% /sys/fs/cgroup
/dev/sdb        9.8G  303M  9.5G  4% /data
/dev/sdal       95G   6.2G  89G  7% /hooks
shm             64M     0   64M  0% /dev/shm
/dev/sdc        2.0G  108K  1.9G  1% /var/log/mongodb-mms-automation
tmpfs          15G   12K   15G  1% /run/secrets/kubernetes.io/serviceaccount
tmpfs          15G     0   15G  0% /proc/acpi
tmpfs          15G     0   15G  0% /proc/scsi
tmpfs          15G     0   15G  0% /sys/firmware
```

存储卷调整之前

```
I have no name!@mongodb-specify-pod-resources-0:/ $ df -h
Filesystem      Size  Used Avail Use% Mounted on
overlay         95G   6.2G  89G  7% /
tmpfs          64M     0   64M  0% /dev
tmpfs          15G     0   15G  0% /sys/fs/cgroup
/dev/sdc        28G  601M  28G  3% /data
/dev/sdal       95G   6.2G  89G  7% /hooks
shm             64M     0   64M  0% /dev/shm
/dev/sdb        2.0G  112K  1.9G  1% /var/log/mongodb-mms-automation
tmpfs          15G   12K   15G  1% /run/secrets/kubernetes.io/serviceaccount
tmpfs          15G     0   15G  0% /proc/acpi
tmpfs          15G     0   15G  0% /proc/scsi
tmpfs          15G     0   15G  0% /sys/firmware
```

存储卷调整之后

3.3、调整MongoDB 副本数量

通过更改MongoDB Operator CRD yaml文件中的member数量，可以实现replicaset数量的自动调整。

```
# 调整前的配置
apiVersion: mongodbcommunity.mongodb.com/v1
kind: MongoDBCommunity
metadata:
  name: mongodb-specify-pod-resources
spec:
  members: 3
  type: ReplicaSet
  version: "4.4.0"
  security:
    authentication:
      modes: ["SCRAM"]
```

```
admin_zzcao_altostrat_com@kubectl:~/mongodb-kubernetes-operator/config/samples$ kubectl get pod -n mongodb
NAME                               READY   STATUS    RESTARTS   AGE
mongodb-kubernetes-operator-774b59c94f-nn5xl   1/1     Running   0          4h24m
mongodb-specify-pod-resources-0            2/2     Running   0          7m39s
mongodb-specify-pod-resources-1            2/2     Running   0          19m
mongodb-specify-pod-resources-2            2/2     Running   0          20m
```

副本数量调整前

```
# 调整后的配置
apiVersion: mongodbcommunity.mongodb.com/v1
kind: MongoDBCommunity
metadata:
  name: mongodb-specify-pod-resources
spec:
  members: 5
  type: ReplicaSet
  version: "4.4.0"
  security:
    authentication:
      modes: ["SCRAM"]
```

```
admin_zzcao_altostrat_com@kubectl:~/mongodb-kubernetes-operator/config/samples$ kubectl get pod -n mongodb
NAME                               READY   STATUS    RESTARTS   AGE
mongodb-kubernetes-operator-774b59c94f-nn5xl   1/1     Running   0          4h30m
mongodb-specify-pod-resources-0            2/2     Running   0          13m
mongodb-specify-pod-resources-1            2/2     Running   0          25m
mongodb-specify-pod-resources-2            2/2     Running   0          25m
mongodb-specify-pod-resources-3            2/2     Running   0          3m17s
mongodb-specify-pod-resources-4            2/2     Running   0          46s
```

副本数量调整后

5.MongoDB备份和恢复

在比较过诸多MongoDB备份和恢复技术后，本文档建议采纳基于GKE的备份和恢复的底层技术，使用[Velero](#)开源工具实现MongoDB的备份和恢复。

4.1. Velero的安装和配置

4.1.1. 安装velero client

```
wget https://github.com/vmware-tanzu/velero/releases/download/v1.8.1/velero-v1.8.1-linux-arm.tar.gz
#解压软件包
tar zxvf velero-v1.8.1-linux-amd64.tar.gz
#将二进制文件移动至/usr/local/bin
sudo cp velero /usr/bin/velero
#验证软件包的安装
velero -h
```

4.1.2. 安装velero server

前提条件

- 创建GCS Bucket，velero server会将备份以及恢复的信息存储在GCS Bucket上。
- 启用workload identity，velero server 需要通过workload identity或者GCP service account key获取对GCS Bucket的操作权限。

详细操作请参考<https://github.com/vmware-tanzu/velero-plugin-for-gcp#setup>。

```
velero install \
--provider gcp \
--plugins velero/velero-plugin-for-gcp:v1.4.0 \
--bucket charles-mongodb-backup-test \
--no-secret \
--sa-annotations
iam.gke.io/gcp-service-account=velero@mongodb-on-gke.iam.gserviceaccount.com \
--backup-location-config
serviceAccount=velero@mongodb-on-gke.iam.gserviceaccount.com \
```

在安装velero server时，velero client会通过Kubernetes API与Kubernetes集群进行交互，因此若环境中存在多个Kubernetes集群时，确保先将kubectl的context切换至目标Kubernetes集群。完成velero server的安装后，velero server会以pod的形式运行在Kubernetes的velero namespace下。

```
admin_zzcao_5t0strat_com@kubectl:~$ kubectl get pod -n velero
NAME          READY   STATUS    RESTARTS   AGE
velero-55bd66b8cb-tc7z4  1/1     Running   0          38m
```

4.2. 基于Velero的备份和恢复的操作

4.2.1 创建velero备份任务

Velero 备份任务可通过 Velero client创建，velero client会与运行在kubernetes集群之上的velero server进行交互，velero server会将备份相关数据上传至GCS Bucket之上。velero会备份两种信息，etcdata数据以及pvc snapshot。

下面命令会将mongodb namespace上所有的资源都进行备份。

```
velero backup create mongodb-XXXX --include-namespaces mongodb
```

```
admin_zxcao_altostrat_com@kubectl:~$ velero backup describe mongodb
Name:          mongodb
Namespace:     velero
Labels:         velero.io/storage-location=default
Annotations:   velero.io/source-cluster-k8s-gitversion=v1.21.6-gke.1500
                velero.io/source-cluster-k8s-major-version=1
                velero.io/source-cluster-k8s-minor-version=21

Phase:        Completed
Errors:       0
Warnings:     0

Namespaces:
  Included:   mongodb
  Excluded:   <none>

Resources:
  Included:    *
  Excluded:   <none>
  Cluster-scoped: auto

Label selector:  <none>
Storage Location: default
Velero-Native Snapshot PVs: auto
TTL: 720h0m0s
Hooks:  <none>
Backup Format Version: 1.1.0
Started: 2022-03-07 06:46:31 +0000 UTC
Completed: 2022-03-07 06:46:40 +0000 UTC
Expiration: 2022-04-06 06:46:31 +0000 UTC
Total items to be backed up: 46
Items backed up:           46
Velero-Native Snapshots: 6 of 6 snapshots completed successfully (specify --details for more information)
```

[查看velero备份计划](#)

GCS存储桶上查看备份数据信息

4.2.2 创建velero恢复任务

```
velero restore create --from=backup mongodb-XX
```

```
admin_zzcao_altostrat_com@kubectl:~$ velero restore describe mongodb-20220307065407
Name:      mongodb-20220307065407
Namespace: velero
Labels:    <none>
Annotations: <none>

Phase:     InProgress
Estimated total items to be restored: 45
Items restored so far: 2

Started:   2022-03-07 06:54:07 +0000 UTC
Completed: <n/a>

Backup:    mongodb

Namespaces:
  Included:  all namespaces found in the backup
  Excluded: <none>

Resources:
  Included:  *
  Excluded:  nodes, events, events.events.k8s.io, backups.velero.io, restores.velero.io, resticrepositories.velero.io
  Cluster-scoped: auto

Namespace mappings: <none>

Label selector: <none>

Restore PVs:  auto

Preserve Service NodePorts: auto
```

查看velero恢复进度

```
admin_zzcao_altostrat_com@kubectl:~$ kubectl get all -n mongodb
NAME                                         READY   STATUS    RESTARTS   AGE
pod/mongodb-kubernetes-operator-774b59c94f-k7f7v   1/1    Running   0          52m
pod/mongodb-second-cluster-0                   2/2    Running   0          48m
pod/mongodb-second-cluster-1                   2/2    Running   0          52m
pod/mongodb-second-cluster-2                   2/2    Running   0          52m

NAME                           TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
service/mongodb-second-cluster-svc  ClusterIP  None        <none>       27017/TCP  52m

NAME                               READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/mongodb-kubernetes-operator  1/1     1           1          52m

NAME                           DESIRED   CURRENT   READY   AGE
replicaset.apps/mongodb-kubernetes-operator-774b59c94f  1         1         1          52m

NAME                               READY   AGE
statefulset.apps/mongodb-second-cluster  3/3     52m
```

Kubernetes成功恢复资源

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	AGE
data-volume-mongodb-second-cluster-0	Bound	pvc-20c2fb73-6e97-4d68-8ada-f9bcefef6ca	501Gi	RWO	standard-rwo	53m
data-volume-mongodb-second-cluster-1	Bound	pvc-37144e82-10b6-442a-840f-62b0bae76c21	501Gi	RWO	standard-rwo	53m
data-volume-mongodb-second-cluster-2	Bound	pvc-4091dc64-dd47-4117-8496-e92385da3c5a	501Gi	RWO	standard-rwo	53m
logs-volume-mongodb-second-cluster-0	Bound	pvc-dc1b5e62-974b-47b2-82df-9f572f562d94	48Gi	RWO	standard-rwo	53m
logs-volume-mongodb-second-cluster-1	Bound	pvc-2c19bb3b-5502-4f49-a5eb-0c12ec65e4d7	48Gi	RWO	standard-rwo	53m
logs-volume-mongodb-second-cluster-2	Bound	pvc-25e5ec05-e1f5-4549-bbb7-00bb2a1a7753	48Gi	RWO	standard-rwo	53m

Kubernetes成功恢复PVC

```
mongodb-second-cluster:SECONDARY> db.stats(1024*1024*1024)
{
  "db" : "workload",
  "collections" : 3,
  "views" : 0,
  "objects" : 88516287,
  "avgObjSize" : 2797.9441070658554,
  "dataSize" : 230.65472356136888,
  "storageSize" : 133.48491668701172,
  "indexes" : 3,
  "indexSize" : 2.5137100219726562,
  "totalSize" : 135.99862670898438,
  "scaleFactor" : 1073741824,
  "fsUsedSize" : 137.09616088867188,
  "fsTotalSize" : 493.07440185546875,
  "ok" : 1,
  "$clusterTime" : {
    "clusterTime" : Timestamp(1646636720, 1),
    "signature" : {
      "hash" : BinData(0,"tsFaF5XfrRhqjWvQ1K6DYKWS3VE="),
      "keyId" : NumberLong("7070698346079322116")
    }
  },
  "operationTime" : Timestamp(1646636720, 1)
}
```

登录MongoDB查看数据恢复状态

4.3. 备份和恢复的基准时间

通过自动化脚本向MongoDB集群内分别注入100G与300G数据，集群的备份与恢复时间记录如下。

MongoDB容量	备份创建时间	集群恢复时间
100GB	5s	6 mins
300GB	5s	8 mins

4.4. 备份和恢复的测试用例

测试备份和恢复前，可以参考[MongoDB造数步骤](#)在MongoDB的部署中打入数据，并准备相应存储容量。

6.MongoDB日志与监控

可通过集成Google Cloud Managed Prometheus服务实现对MongoDB集群的监控。

5.1、安装Managed Prometheus Collector, Collector以Pod方式运行，是精简版的Prometheus。

Collector收集exporter提供的指标，并上传至Managed Prometehus云端。

```
kubectl apply -f  
https://raw.githubusercontent.com/GoogleCloudPlatform/prometheus-engine/v0.3.0/manifests/setup.yaml
```

```
kubectl apply -f  
https://raw.githubusercontent.com/GoogleCloudPlatform/prometheus-engine/v0.3.1/manifests/operator.yaml
```

5.2、安装MongoDB Exporter, MongoDB Exporter用于收集MongoDB集群指标，并将指标提供给Managed Prometheus Collector。

```
helm install mongodb-exporter prometheus-community/prometheus-mongodb-exporter -f value.yaml -n mongodb
```

Value.yaml内容如下，将<mongodb-url>替换成[3.5步骤的输出](#)。

```
mongodb:  
  uri: "<mongodb-url>"  
  serviceMonitor:  
    enabled: false
```

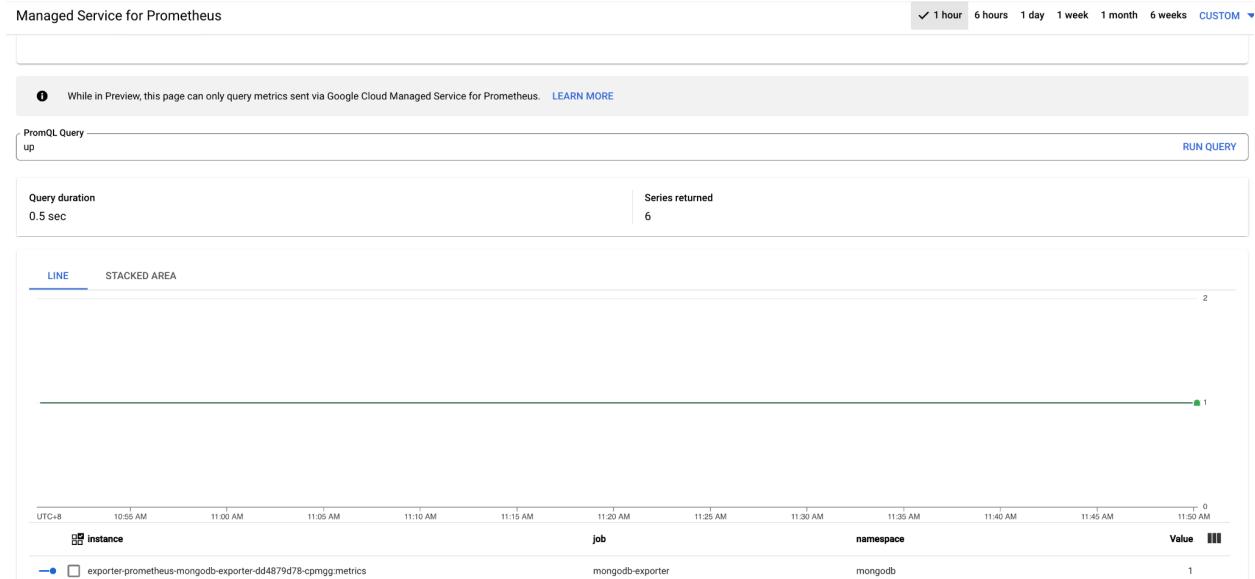
5.3、创建Pod Monitoring资源，PodMonitoring用于定义metric target，Prometheus collector通过pod monitoring定义收集相应endpoint的指标。

```
kubectl apply -f mongodb-exporter-pm.yaml
```

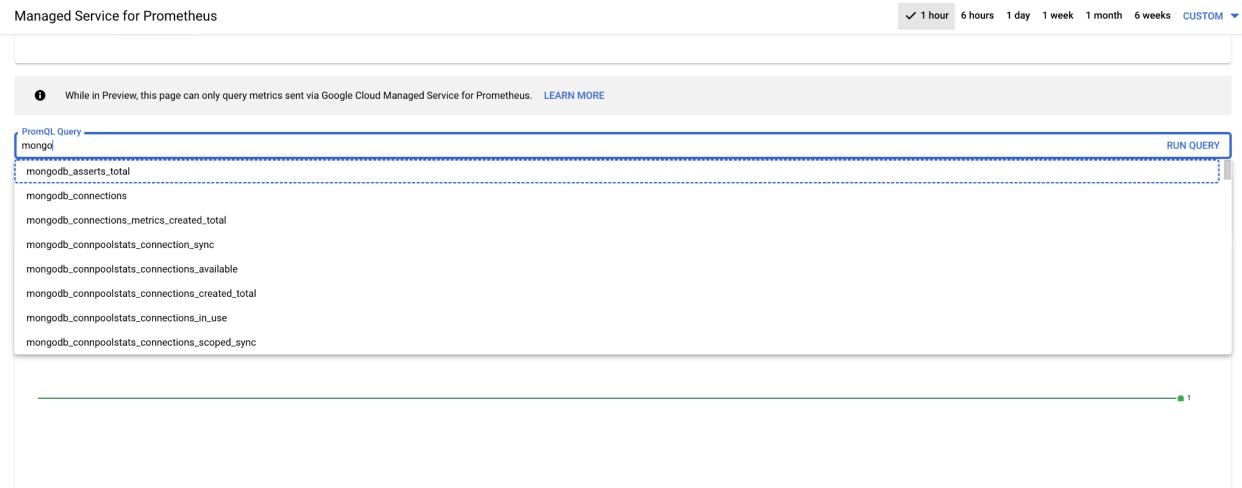
Mongodb-exporter-pm.yaml 内容如下：

```
apiVersion: monitoring.googleapis.com/v1alpha1  
kind: PodMonitoring  
metadata:  
  name: prom-example  
spec:  
  selector:  
    matchLabels:  
      app.kubernetes.io/instance: mongodb-exporter  
  endpoints:  
  - port: metrics  
    interval: 30s
```

5.4、到Google Cloud Prometheus控制台上查看指标



确认exporter实例已被Prometheus服务发现



查看mongodb相关指标

5.5、除了Prometheus控制台，也可以通过cloud monitoring的方式查看MongoDB指标

The screenshot shows the Google Cloud Platform Metrics Explorer interface. On the left sidebar, under the 'Metrics explorer' section, there is a dropdown menu with the option 'Metrics Scope' selected, showing '1 project'. The main area displays a 'SELECT A METRIC' dropdown menu. The search bar at the top contains the text 'pro'. Below the search bar, the results are categorized into 'POPULAR RESOURCES' and 'ACTIVE RESOURCES'. Under 'ACTIVE RESOURCES', there is a section for 'Prometheus Target' which lists 473 metrics. One metric is highlighted: 'MongoDB.asserts.total'. To the right of the dropdown, there is a placeholder text 'Select a metric to start' and a timeline from 11:15 AM to 11:45 AM.

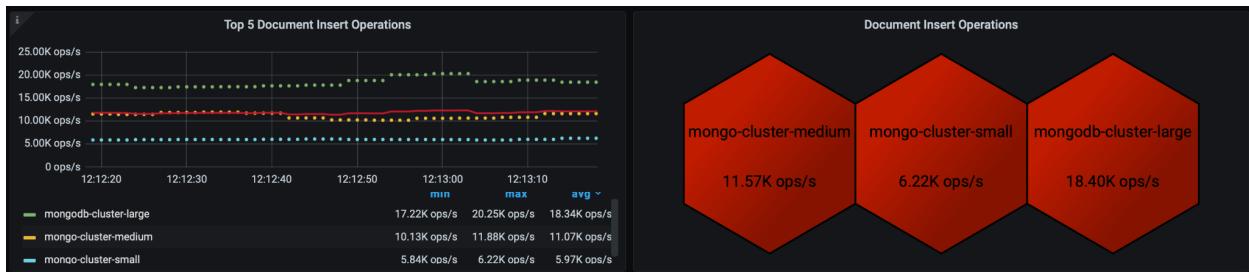
通过cloud monitoring的方式查看MongoDB指标

7.MongoDB性能基准测试

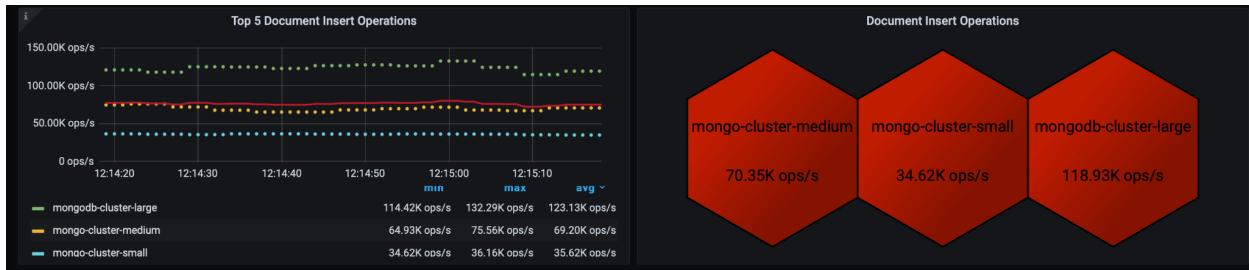
MongoDB性能基准测试参考(Percona)

本次测试采用开源工具POCDriver对不同规模的MongoDB集群进行性能测试，每个集群分别测试了单个batch插入1条文档，单个batch插入10条文档，单个batch插入100条文档，性能测试结果记录如。

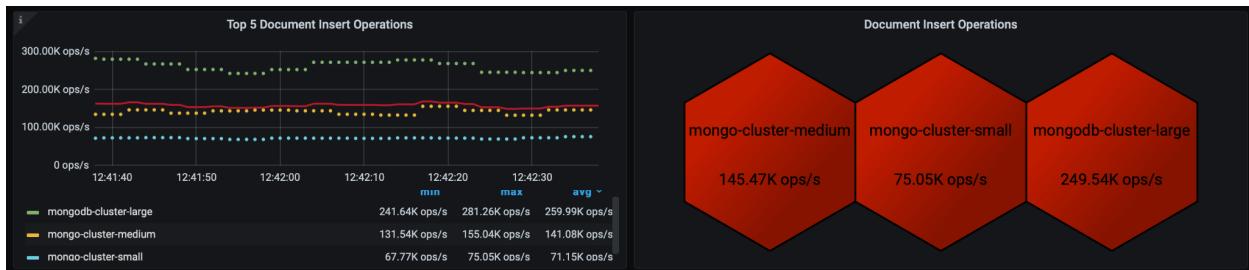
集群规模	集群配置	写QPS(batch:1)	写QPS(batch:10)	写QPS(batch:100)
small	2C/8GB /512GB SSD	6.22k	34.62k	75.05k
medium	4C/16GB/1TB SSD	11.57k	70.35k	145.47k
large	8C/32G/2TB SSD	18.40k	118.93k	249.54k



MongoDB集群Qps(1document/batch)



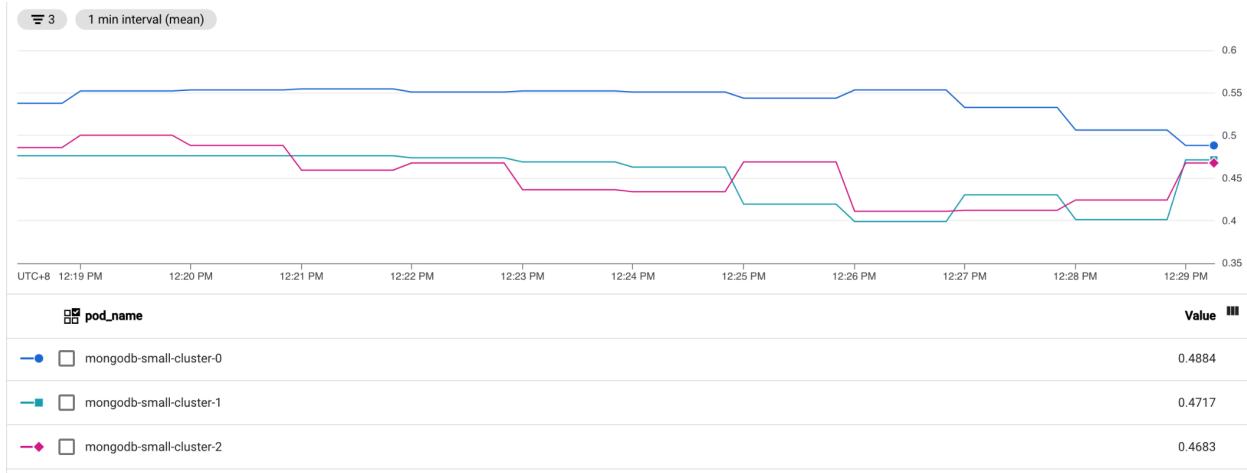
MongoDB集群Qps(10documents/batch)



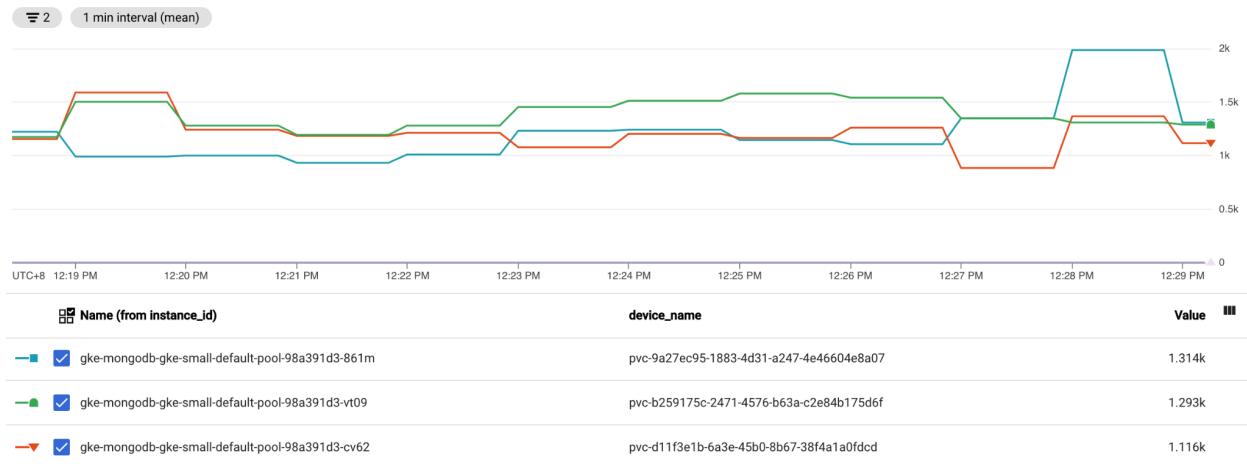
MongoDB集群Qps(100documents/batch)

容器与虚拟机资源占用情况





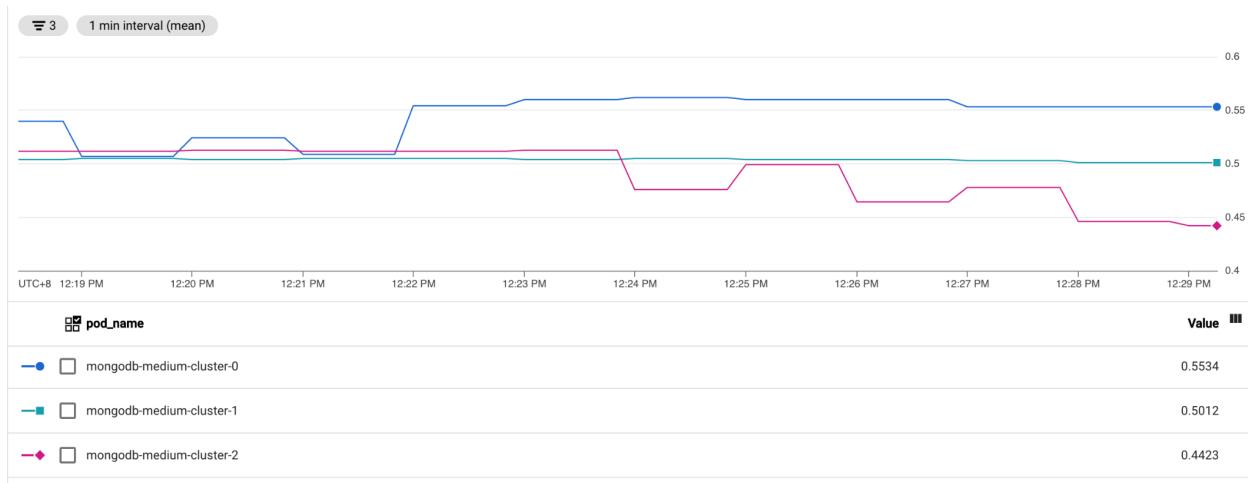
内存利用率 (small)



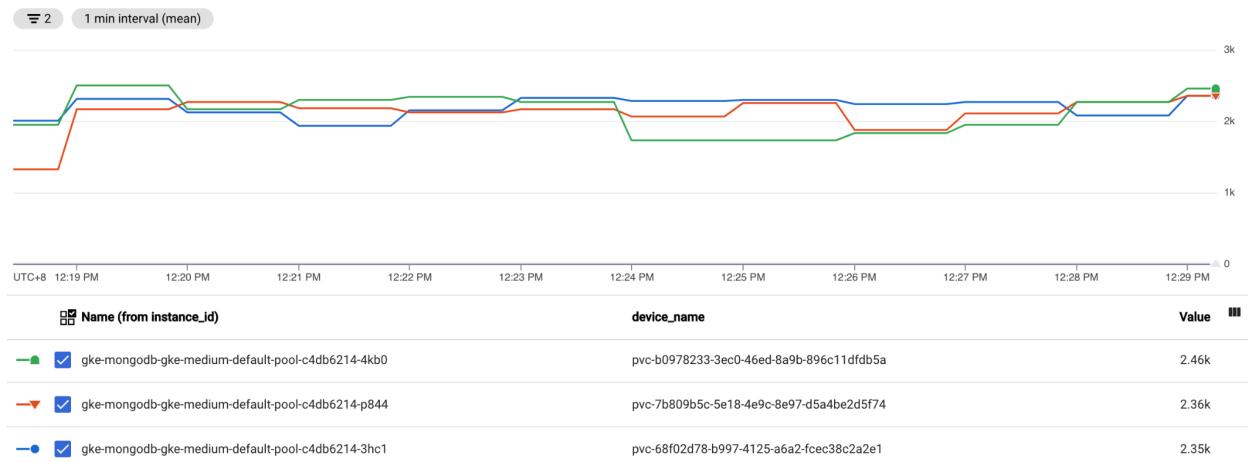
存储IOPS (Small)



CPU利用率 (medium)



内存利用率 (medium)



存储IOPS (medium)



CPU利用率(large)



内存利用率(large)



存储IOPS (large)

MongoDB性能测试方法请参考：

- 使用基准测试工具sysbench的[文档](#)
- 使用基准测试工具POCDriver的[文档](#)

1. 访问MongoDB on GKE

7.1. 从内部访问

7.2. 从外部访问

<https://appfleet.com/blog/orchaestrating-mongodb-over-kubernetes/>

8.附录:

8.1 参考的MongoDBCommunity配置文件

```
---
apiVersion: mongodbcommunity.mongodb.com/v1
kind: MongoDBCommunity
metadata:
  name: example-mongodb
spec:
  members: 3
  type: ReplicaSet
  version: "4.4.13"
  security:
    authentication:
      modes: ["SCRAM"]
  users:
    - name: my-user
      db: admin
      passwordSecretRef: # a reference to the secret that will be used to generate
the user's password
        name: my-user-password
      roles:
        - name: clusterAdmin
          db: admin
        - name: userAdminAnyDatabase
          db: admin
    scramCredentialsSecretName: my-scram
statefulSet:
  spec:
    template:
      spec:
        # resources can be specified by applying an override
        # per container name.
        containers:
```

```

- name: mongod
  resources:
    limits:
      cpu: "2"
      memory: 2000M
    requests:
      cpu: "0.2"
      memory: 200M
- name: mongodb-agent
  resources:
    limits:
      cpu: "0.2"
      memory: 250M
    requests:
      cpu: "0.2"
      memory: 200M
# Specifies a size for the data volume different from the default 10Gi
volumeClaimTemplates:
- metadata:
    name: data-volume
  spec:
    accessModes: [ "ReadWriteOnce" ]
    resources:
      requests:
        storage: 15Gi
additionalMongodConfig:
  storage.wiredTiger.engineConfig.journalCompressor: zlib

# the user credentials will be generated from this secret
# once the credentials are generated, this secret is no longer required
---
apiVersion: v1
kind: Secret
metadata:
  name: my-user-password
type: Opaque
stringData:
  password: <your-password-here>

```

8.2 从vpc内部访问mongo cluster

- 获取连接串

```
k get secret example-mongodb-admin-my-user \
-o json | jq -r '.data | with_entries(.value |= @base64d)
```

```
{
  "connectionString.standard": "mongodb://my-user:changeme@example-mongodb-0.example-mongodb-svc.default.svc.goodvm.local:27017,example-mongodb-1.example-mongodb-svc.default.svc.goodvm.local:27017,example-mongodb-2.example-mongodb-svc.default.svc.goodvm.local:27017/admin?ssl=false",
  "connectionString.standardSrv": "mongodb+srv://my-user:changeme@example-mongodb-svc.default.svc.goodvm.local/admin?ssl=false",
  "password": "changeme",
  "username": "my-user"
}
```

- Mongosh 连接后创建管理员用户

```
mongosh
"mongodb://my-user:changeme@example-mongodb-0.example-mongodb-svc.default.svc.goodvm.local:27017,example-mongodb-1.example-mongodb-svc.default.svc.goodvm.local:27017,example-mongodb-2.example-mongodb-svc.default.svc.goodvm.local:27017/admin?ssl=false"
```

```
example-mongodb [primary] admin> db.createUser(
  {
    user: 'admin',
    pwd: 'password',
    roles: [ { role: 'root', db: 'admin' } ]
  }
);
```

8.3 velero workload identity模式安装参考

```
kubectl create namespace velero
kubectl create serviceaccount velero --namespace velero
gcloud iam service-accounts add-iam-policy-binding
cliu201-sa@cliu201.iam.gserviceaccount.com --role roles/iam.workloadIdentityUser
--member "serviceAccount:cliu201.svc.id.goog[velero/velero]"
```

```
velero install \n
--provider gcp \n
--plugins velero/velero-plugin-for-gcp:v1.4.0 \n
--bucket cliu201-us \n
--no-secret \n
--sa-annotations
iam.gke.io/gcp-service-account=cliu201-sa@cliu201.iam.gserviceaccount.com \n
--backup-location-config serviceAccount=cliu201-sa@cliu201.iam.gserviceaccount.com

# 手动添加service account annotation
kubectl annotate serviceaccount velero
iam.gke.io/gcp-service-account=cliu201-sa@cliu201.iam.gserviceaccount.com
--namespace velero
```

8.4 mongodb-exporter service和collector的关联关系建立

```
> k describe svc mongodb-exporter-prometheus-mongodb-exporter
Name:           mongodb-exporter-prometheus-mongodb-exporter
Namespace:      default
Labels:         app.kubernetes.io/instance=mongodb-exporter
                app.kubernetes.io/managed-by=Helm
                app.kubernetes.io/name=prometheus-mongodb-exporter
                helm.sh/chart=prometheus-mongodb-exporter-2.9.0
Annotations:   cloud.google.com/neg: {"ingress":true}
                meta.helm.sh/release-name: mongodb-exporter
                meta.helm.sh/release-namespace: default
Selector:       app.kubernetes.io/instance=mongodb-exporter,app.kubernetes.io/name=prometheus-mongodb-exporter
Type:          ClusterIP
IP Family Policy: SingleStack
IP Families:   IPv4
IP:            10.44.13.5
IPs:           10.44.13.5
Port:          metrics 9216/TCP
TargetPort:    metrics/TCP
Endpoints:     10.40.2.5:9216
Session Affinity: None
Events:        <none>
� ↵ ~/k8s/0323
```

```
! mongodb.com_v1_mongodbcommunity_cr.yaml M      ! mongodb-exporter-pm.yaml X  ! tt.yaml
0323 > ! mongodb-exporter-pm.yaml > {} spec > {} selector
  1   apiVersion: monitoring.googleapis.com/v1alpha1
  2   kind: PodMonitoring
  3   metadata:
  4     name: prom-example
  5   spec:
  6     selector:
  7       matchLabels:
  8         app.kubernetes.io/instance: mongodb-exporter
  9     endpoints:
10       - port: metrics
11         interval: 30s
```